



CUDA **CUFFT Library**

PG-00000-003_V2.3
June, 2009

Confidential Information

Published by
NVIDIA Corporation
2701 San Tomas Expressway
Santa Clara, CA 95050

Notice

This source code is subject to NVIDIA ownership rights under U.S. and international Copyright laws.

This software and the information contained herein is PROPRIETARY and CONFIDENTIAL to NVIDIA and is being provided under the terms and conditions of a Non-Disclosure Agreement. Any reproduction or disclosure to any third party without the express written consent of NVIDIA is prohibited.

NVIDIA MAKES NO REPRESENTATION ABOUT THE SUITABILITY OF THIS SOURCE CODE FOR ANY PURPOSE. IT IS PROVIDED "AS IS" WITHOUT EXPRESS OR IMPLIED WARRANTY OF ANY KIND. NVIDIA DISCLAIMS ALL WARRANTIES WITH REGARD TO THIS SOURCE CODE, INCLUDING ALL IMPLIED WARRANTIES OF MERCHANTABILITY, NONINFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE. IN NO EVENT SHALL NVIDIA BE LIABLE FOR ANY SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, OR ANY DAMAGES WHATSOEVER RESULTING FROM LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR OTHER TORTIOUS ACTION, ARISING OUT OF OR IN CONNECTION WITH THE USE OR PERFORMANCE OF THIS SOURCE CODE.

U.S. Government End Users. This source code is a "commercial item" as that term is defined at 48 C.F.R. 2.101 (OCT 1995), consisting of "commercial computer software" and "commercial computer software documentation" as such terms are used in 48 C.F.R. 12.212 (SEPT 1995) and is provided to the U.S. Government only as a commercial end item. Consistent with 48 C.F.R.12.212 and 48 C.F.R. 227.7202-1 through 227.7202-4 (JUNE 1995), all U.S. Government End Users acquire the source code with only those rights set forth herein.

Trademarks

NVIDIA, CUDA, and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the United States and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

© 2006–2009 by NVIDIA Corporation. All rights reserved.

Table of Contents

CUFFT Library	1
CUFFT Types and Definitions	1
Type cufftHandle	2
Type cufftResult	2
Type cufftReal	2
Type cufftDoubleReal	3
Type cufftComplex	3
Type cufftDoubleComplex	3
CUFFT Transform Types	3
CUFFT Transform Directions	4
CUFFT API Functions	5
Function cufftPlan1d()	6
Function cufftPlan2d()	7
Function cufftPlan3d()	7
Function cufftDestroy()	8
Function cufftExecC2C()	8
Function cufftExecR2C()	9
Function cufftExecC2R()	10
Function cufftExecZ2Z()	11
Function cufftExecD2Z()	12
Function cufftExecZ2D()	12
Accuracy and Performance	13
CUFFT Code Examples	15
1D Complex-to-Complex Transforms	15
1D Real-to-Complex Transforms	16
2D Complex-to-Complex Transforms	16
2D Complex-to-Real Transforms	17
3D Complex-to-Complex Transforms	18

CUFFT Library

This document describes CUFFT, the NVIDIA[®] CUDA[™] (compute unified device architecture) Fast Fourier Transform (FFT) library. The FFT is a divide-and-conquer algorithm for efficiently computing discrete Fourier transforms of complex or real-valued data sets, and it is one of the most important and widely used numerical algorithms, with applications that include computational physics and general signal processing. The CUFFT library provides a simple interface for computing parallel FFTs on an NVIDIA GPU, which allows users to leverage the floating-point power and parallelism of the GPU without having to develop a custom, GPU-based FFT implementation.

FFT libraries typically vary in terms of supported transform sizes and data types. For example, some libraries only implement Radix-2 FFTs, restricting the transform size to a power of two, while other implementations support arbitrary transform sizes. This version of the CUFFT library supports the following features:

- ❑ 1D, 2D, and 3D transforms of complex and real-valued data
- ❑ Batch execution for doing multiple 1D transforms in parallel
- ❑ 2D and 3D transform sizes in the range [2, 16384] in any dimension
- ❑ 1D transform sizes up to 8 million elements
- ❑ In-place and out-of-place transforms for real and complex data
- ❑ Double-precision transforms on compatible hardware (GT200 and later GPUs)

CUFFT Types and Definitions

The next sections describe the CUFFT types and transform directions:

- ❑ [“Type cufftHandle” on page 2](#)
- ❑ [“Type cufftResult” on page 2](#)
- ❑ [“Type cufftReal” on page 2](#)

- ❑ “Type `cufftDoubleReal`” on page 3
- ❑ “Type `cufftComplex`” on page 3
- ❑ “Type `cufftDoubleComplex`” on page 3
- ❑ “CUFFT Transform Types” on page 3
- ❑ “CUFFT Transform Directions” on page 4

Type `cufftHandle`

```
typedef unsigned int cufftHandle;
```

is a handle type used to store and access CUFFT plans. For example, the user receives a handle after creating a CUFFT plan and uses this handle to execute the plan.

Type `cufftResult`

```
typedef enum cufftResult_t cufftResult;
```

is an enumeration of values used exclusively as API function return values. The possible return values are defined as follows:

Return Values

CUFFT_SUCCESS	Any CUFFT operation is successful.
CUFFT_INVALID_PLAN	CUFFT is passed an invalid plan handle.
CUFFT_ALLOC_FAILED	CUFFT failed to allocate GPU memory.
CUFFT_INVALID_TYPE	The user requests an unsupported type.
CUFFT_INVALID_VALUE	The user specifies a bad memory pointer.
CUFFT_INTERNAL_ERROR	Used for all internal driver errors.
CUFFT_EXEC_FAILED	CUFFT failed to execute an FFT on the GPU.
CUFFT_SETUP_FAILED	The CUFFT library failed to initialize.
CUFFT_SHUTDOWN_FAILED	The CUFFT library failed to shut down.
CUFFT_INVALID_SIZE	The user specifies an unsupported FFT size.

Type `cufftReal`

```
typedef float cufftReal;
```

is a single-precision, floating-point real data type.

Type `cufftDoubleReal`

```
typedef double cufftDoubleReal;
```

is a double-precision, floating-point real data type.

Type `cufftComplex`

```
typedef cuComplex cufftComplex;
```

is a single-precision, floating-point complex data type that consists of interleaved real and imaginary components.

Type `cufftDoubleComplex`

```
typedef cuDoubleComplex cufftDoubleComplex;
```

is a double-precision, floating-point complex data type that consists of interleaved real and imaginary components.

CUFFT Transform Types

The CUFFT library supports complex- and real-data transforms. The **`cufftType`** data type is an enumeration of the types of transform data supported by CUFFT:

```
typedef enum cufftType_t {  
    CUFFT_R2C = 0x2a, // Real to complex (interleaved)  
    CUFFT_C2R = 0x2c, // Complex (interleaved) to real  
    CUFFT_C2C = 0x29, // Complex to complex, interleaved  
    CUFFT_D2Z = 0x6a, // Double to double-complex  
    CUFFT_Z2D = 0x6c, // Double-complex to double  
    CUFFT_Z2Z = 0x69  // Double-complex to double-complex  
} cufftType;
```

For complex FFTs, the input and output arrays must interleave the real and imaginary parts (the **`cufftComplex`** type). The transform size in each dimension is the number of **`cufftComplex`** elements. The **`CUFFT_C2C`** constant can be passed to any plan creation function to configure a single-precision complex-to-complex FFT. Pass the **`CUFFT_Z2Z`** constant to configure a double-precision complex-to-complex FFT.

For real-to-complex FFTs, the output array holds only the non-redundant complex coefficients. So for an N -element transform, the output array holds $N/2+1$ `cufftComplex` terms. For higher-dimensional real transforms of the form $N_0 \times N_1 \times \dots \times N_n$, the last dimension is cut in half such that the output data is $N_0 \times N_1 \times \dots \times (N_n/2+1)$ complex elements. Therefore, in order to perform an in-place FFT, the user has to pad the input array in the last dimension to $(N_n/2+1)$ complex elements or $2 \times (N/2+1)$ real elements. Note that the real-to-complex transform is implicitly forward. Passing the `CUFFT_R2C` constant to any plan creation function configures a single-precision real-to-complex FFT. Passing the `CUFFT_D2Z` constant configures a double-precision real-to-complex FFT.

The requirements for complex-to-real FFTs are similar to those for real-to-complex. In this case, the input array holds only the non-redundant, $N/2+1$ complex coefficients from a real-to-complex transform. The output is simply N elements of type `cufftReal`. However, for an in-place transform, the input size must be padded to $2 \times (N/2+1)$ real elements. The complex-to-real transform is implicitly inverse. Passing the `CUFFT_C2R` constant to any plan creation function configures a single-precision complex-to-real FFT. Passing `CUFFT_Z2D` constant configures a double-precision complex-to-real FFT.

For 1D complex-to-complex transforms, the stride between signals in a batch is assumed to be the number of `cufftComplex` elements in the logical transform size. However, for real-data FFTs, the distance between signals in a batch depends on whether the transform is in-place or out-of-place. For in-place FFTs, the input stride is assumed to be $2 \times (N/2+1)$ `cufftReal` elements or $N/2+1$ `cufftComplex` elements. For out-of-place transforms, the input and output strides match the logical transform size (N) and the non-redundant size ($N/2+1$), respectively.

CUFFT Transform Directions

The CUFFT library defines forward and inverse Fast Fourier Transforms according to the sign of the complex exponential term:

```
#define CUFFT_FORWARD -1
#define CUFFT_INVERSE 1
```

For higher-dimensional transforms (2D and 3D), CUFFT performs FFTs in row-major or C order. For example, if the user requests a 3D transform plan for sizes X , Y , and Z , CUFFT transforms along Z , Y , and then X . The user can configure column-major FFTs by simply changing the order of the size parameters to the plan creation API functions.

CUFFT performs un-normalized FFTs; that is, performing a forward FFT on an input data set followed by an inverse FFT on the resulting set yields data that is equal to the input scaled by the number of elements. Scaling either transform by the reciprocal of the size of the data set is left for the user to perform as seen fit.

CUFFT API Functions

The CUFFT API is modeled after FFTW (see <http://www.fftw.org>), which is one of the most popular and efficient CPU-based FFT libraries. FFTW provides a simple configuration mechanism called a *plan* that completely specifies the optimal—that is, the minimum floating-point operation (flop)—plan of execution for a particular FFT size and data type. The advantage of this approach is that once the user creates a plan, the library stores whatever state is needed to execute the plan multiple times without recalculation of the configuration. The FFTW model works well for CUFFT because different kinds of FFTs require different thread configurations and GPU resources, and plans are a simple way to store and reuse configurations.

The CUFFT library initializes internal data upon the first invocation of an API function. Therefore, all API functions could return the **CUFFT_SETUP_FAILED** error code if the library fails to initialize. CUFFT shuts down automatically when all user-created FFT plans are destroyed.

The CUFFT functions are as follows:

- ❑ “Function `cufftPlan1d()`” on page 6
- ❑ “Function `cufftPlan2d()`” on page 7
- ❑ “Function `cufftPlan3d()`” on page 7
- ❑ “Function `cufftDestroy()`” on page 8

- ❑ “Function `cufftExecC2C()`” on page 8
- ❑ “Function `cufftExecR2C()`” on page 9
- ❑ “Function `cufftExecC2R()`” on page 10
- ❑ “Function `cufftExecZ2Z()`” on page 11
- ❑ “Function `cufftExecD2Z()`” on page 12
- ❑ “Function `cufftExecZ2D()`” on page 12

Function `cufftPlan1d()`

```
cufftResult
cufftPlan1d( cufftHandle *plan, int nx, cufftType type,
             int batch );
```

creates a 1D FFT plan configuration for a specified signal size and data type. The `batch` input parameter tells CUFFT how many 1D transforms to configure.

Input

<code>plan</code>	Pointer to a <code>cufftHandle</code> object
<code>nx</code>	The transform size (e.g., 256 for a 256-point FFT)
<code>type</code>	The transform data type (e.g., <code>CUFFT_C2C</code> for complex to complex)
<code>batch</code>	Number of transforms of size <code>nx</code>

Output

<code>plan</code>	Contains a CUFFT 1D plan handle value
-------------------	---------------------------------------

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_SIZE	The <code>nx</code> parameter is not a supported size.
CUFFT_INVALID_TYPE	The <code>type</code> parameter is not supported.
CUFFT_ALLOC_FAILED	Allocation of GPU resources for the plan failed.
CUFFT_SUCCESS	CUFFT successfully created the FFT plan.

Function `cufftPlan2d()`

```
cufftResult
cufftPlan2d( cufftHandle *plan, int nx, int ny,
             cufftType type );
```

creates a 2D FFT plan configuration according to specified signal sizes and data type. This function is the same as `cufftPlan1d()` except that it takes a second size parameter, `ny`, and does not support batching.

Input

<code>plan</code>	Pointer to a <code>cufftHandle</code> object
<code>nx</code>	The transform size in the <i>X</i> dimension (number of rows)
<code>ny</code>	The transform size in the <i>Y</i> dimension (number of columns)
<code>type</code>	The transform data type (e.g., <code>CUFFT_C2R</code> for complex to real)

Output

<code>plan</code>	Contains a CUFFT 2D plan handle value
-------------------	---------------------------------------

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_SIZE	The <code>nx</code> or <code>ny</code> parameter is not a supported size.
CUFFT_INVALID_TYPE	The <code>type</code> parameter is not supported.
CUFFT_ALLOC_FAILED	Allocation of GPU resources for the plan failed.
CUFFT_SUCCESS	CUFFT successfully created the FFT plan.

Function `cufftPlan3d()`

```
cufftResult
cufftPlan3d( cufftHandle *plan, int nx, int ny, int nz,
             cufftType type );
```

creates a 3D FFT plan configuration according to specified signal sizes and data type. This function is the same as `cufftPlan2d()` except that it takes a third size parameter `nz`. :

Input

<code>plan</code>	Pointer to a <code>cufftHandle</code> object
<code>nx</code>	The transform size in the <i>X</i> dimension
<code>ny</code>	The transform size in the <i>Y</i> dimension

Input (continued)

<code>nz</code>	The transform size in the Z dimension
<code>type</code>	The transform data type (e.g., <code>CUFFT_R2C</code> for real to complex)

Output

<code>plan</code>	Contains a CUFFT 3D plan handle value
-------------------	---------------------------------------

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_SIZE	Parameter <code>nx</code> , <code>ny</code> , or <code>nz</code> is not a supported size.
CUFFT_INVALID_TYPE	The <code>type</code> parameter is not supported.
CUFFT_ALLOC_FAILED	Allocation of GPU resources for the plan failed.
CUFFT_SUCCESS	CUFFT successfully created the FFT plan.

Function `cufftDestroy()`

cufftResult

cufftDestroy(cufftHandle plan);

frees all GPU resources associated with a CUFFT plan and destroys the internal plan data structure. This function should be called once a plan is no longer needed to avoid wasting GPU memory.

Input

<code>plan</code>	The <code>cufftHandle</code> object of the plan to be destroyed.
-------------------	--

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_SHUTDOWN_FAILED	CUFFT library failed to shut down.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_SUCCESS	CUFFT successfully destroyed the FFT plan.

Function `cufftExecC2C()`

cufftResult

cufftExecC2C(cufftHandle plan, cufftComplex *idata, cufftComplex *odata, int direction);

executes a CUFFT single-precision complex-to-complex transform plan as specified by `direction`. CUFFT uses as input data the GPU

memory pointed to by the `idata` parameter. This function stores the Fourier coefficients in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform.

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
<code>idata</code>	Pointer to the single-precision complex input data (in GPU memory) to transform
<code>odata</code>	Pointer to the single-precision complex output data (in GPU memory)
<code>direction</code>	The transform direction: <code>CUFFT_FORWARD</code> or <code>CUFFT_INVERSE</code>

Output

<code>odata</code>	Contains the complex Fourier coefficients
--------------------	---

Return Values

<code>CUFFT_SETUP_FAILED</code>	CUFFT library failed to initialize.
<code>CUFFT_INVALID_PLAN</code>	The <code>plan</code> parameter is not a valid handle.
<code>CUFFT_INVALID_VALUE</code>	The <code>idata</code> , <code>odata</code> , and/or <code>direction</code> parameter is not valid.
<code>CUFFT_EXEC_FAILED</code>	CUFFT failed to execute the transform on GPU.
<code>CUFFT_SUCCESS</code>	CUFFT successfully executed the FFT plan.

Function `cufftExecR2C()`

`cufftResult`

```
cufftExecR2C( cufftHandle plan, cufftReal *idata,
               cufftComplex *odata );
```

executes a CUFFT single-precision real-to-complex (implicitly forward) transform plan. CUFFT uses as input data the GPU memory pointed to by the `idata` parameter. This function stores the non-redundant Fourier coefficients in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform (See [“CUFFT Transform Types” on page 3](#) for details on real data FFTs.)

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
-------------------	--

Input (continued)

<code>idata</code>	Pointer to the single-precision real input data (in GPU memory) to transform
<code>odata</code>	Pointer to the single-precision complex output data (in GPU memory)

Output

<code>odata</code>	Contains the complex Fourier coefficients
--------------------	---

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_INVALID_VALUE	The <code>idata</code> and/or <code>odata</code> parameter is not valid.
CUFFT_EXEC_FAILED	CUFFT failed to execute the transform on GPU.
CUFFT_SUCCESS	CUFFT successfully executed the FFT plan.

Function `cufftExecC2R()`

cufftResult

```
cufftExecC2R( cufftHandle plan, cufftComplex *idata,
               cufftReal *odata );
```

executes a CUFFT single-precision complex-to-real (implicitly inverse) transform plan. CUFFT uses as input data the GPU memory pointed to by the `idata` parameter. The input array holds only the non-redundant complex Fourier coefficients. This function stores the real output values in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform. (See [“CUFFT Transform Types” on page 3](#) for details on real data FFTs.)

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
<code>idata</code>	Pointer to the single-precision complex input data (in GPU memory) to transform
<code>odata</code>	Pointer to the single-precision real output data (in GPU memory)

Output

<code>odata</code>	Contains the real-valued output data
--------------------	--------------------------------------

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_INVALID_VALUE	The <code>idata</code> and/or <code>odata</code> parameter is not valid.
CUFFT_EXEC_FAILED	CUFFT failed to execute the transform on GPU.
CUFFT_SUCCESS	CUFFT successfully executed the FFT plan.

Function `cufftExecZ2Z()`

```
cufftResult
cufftExecZ2Z( cufftHandle plan,
               cufftDoubleComplex *idata,
               cufftDoubleComplex *odata, int direction );
```

executes a CUFFT double-precision complex-to-complex transform plan as specified by `direction`. CUFFT uses as input data the GPU memory pointed to by the `idata` parameter. This function stores the Fourier coefficients in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform.

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
<code>idata</code>	Pointer to the double-precision complex input data (in GPU memory) to transform
<code>odata</code>	Pointer to the double-precision complex output data (in GPU memory)
<code>direction</code>	The transform direction: <code>CUFFT_FORWARD</code> or <code>CUFFT_INVERSE</code>

Output

<code>odata</code>	Contains the complex Fourier coefficients
--------------------	---

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_INVALID_VALUE	The <code>idata</code> , <code>odata</code> , and/or <code>direction</code> parameter is not valid.
CUFFT_EXEC_FAILED	CUFFT failed to execute the transform on GPU.
CUFFT_SUCCESS	CUFFT successfully executed the FFT plan.

Function cufftExecD2Z()

```
cufftResult
cufftExecD2Z( cufftHandle plan, cufftDoubleReal *idata,
                cufftDoubleComplex *odata );
```

executes a CUFFT double-precision real-to-complex (implicitly forward) transform plan. CUFFT uses as input data the GPU memory pointed to by the `idata` parameter. This function stores the non-redundant Fourier coefficients in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform (See [“CUFFT Transform Types” on page 3](#) for details on real data FFTs.)

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
<code>idata</code>	Pointer to the double-precision real input data (in GPU memory) to transform
<code>odata</code>	Pointer to the double-precision complex output data (in GPU memory)

Output

<code>odata</code>	Contains the complex Fourier coefficients
--------------------	---

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_INVALID_VALUE	The <code>idata</code> and/or <code>odata</code> parameter is not valid.
CUFFT_EXEC_FAILED	CUFFT failed to execute the transform on GPU.
CUFFT_SUCCESS	CUFFT successfully executed the FFT plan.

Function cufftExecZ2D()

```
cufftResult
cufftExecZ2D( cufftHandle plan,
                cufftDoubleComplex *idata,
                cufftDoubleReal *odata );
```

executes a CUFFT double-precision complex-to-real (implicitly inverse) transform plan. CUFFT uses as input data the GPU memory pointed to by the `idata` parameter. The input array holds only the non-redundant complex Fourier coefficients. This function stores the

real output values in the `odata` array. If `idata` and `odata` are the same, this method does an in-place transform. (See “[CUFFT Transform Types](#)” on page 3 for details on real data FFTs.)

Input

<code>plan</code>	The <code>cufftHandle</code> object for the plan to update
<code>idata</code>	Pointer to the double-precision complex input data (in GPU memory) to transform
<code>odata</code>	Pointer to the double-precision real output data (in GPU memory)

Output

<code>odata</code>	Contains the real-valued output data
--------------------	--------------------------------------

Return Values

CUFFT_SETUP_FAILED	CUFFT library failed to initialize.
CUFFT_INVALID_PLAN	The <code>plan</code> parameter is not a valid handle.
CUFFT_INVALID_VALUE	The <code>idata</code> and/or <code>odata</code> parameter is not valid.
CUFFT_EXEC_FAILED	CUFFT failed to execute the transform on GPU.
CUFFT_SUCCESS	CUFFT successfully executed the FFT plan.

Accuracy and Performance

The CUFFT library implements several FFT algorithms, each having different performance and accuracy. The best performance paths correspond to transform sizes that meet two criteria:

1. Fit in CUDA's shared memory
2. Are powers of a single factor (for example, powers of two)

These transforms are also the most accurate due to the numeric stability of the chosen FFT algorithm. For transform sizes that meet the first criterion but not second, CUFFT uses a more general mixed-radix FFT algorithm that is usually slower and less numerically accurate.

Therefore, if possible it is best to use sizes that are powers of two or four, or powers of other small primes (such as, three, five, or seven). In addition, the power-of-two FFT algorithm in CUFFT makes maximum use of shared memory by blocking sub-transforms for signals that do not meet the first criterion.

For transform sizes that do not meet either criteria above, CUFFT uses an out-of-place, mixed-radix algorithm that stores all intermediate results in CUDA's global GPU memory. Although this algorithm uses optimized transform modules for many factors, it has generally lower performance because global memory has less bandwidth than shared memory. The one exception is large 1D transforms, where CUFFT uses a distributed algorithm that performs a 1D FFT using a 2D FFT, where the dimensions of the 2D transform are factors of the 1D size. This path attempts to utilize the faster transforms mentioned above even if the signal size is too large to fit in CUDA's shared memory.

Many FFT algorithms for real data exploit the conjugate symmetry property to reduce computation and memory cost by roughly half. However, CUFFT does not implement any specialized algorithms for real data, and so there is no direct performance benefit to using real-to-complex (or complex-to-real) plans instead of complex-to-complex. For this release, the real data API exists primarily for convenience, so that users do not have to build interleaved complex data from a real data source before using the library. For 1D transforms, the performance for real data will either match or be less than the complex equivalent (due to an extra copy in some cases). However, there is usually a performance benefit to using real data for 2D and 3D FFTs, since all transforms but the last dimension operate on roughly half the logical signal size

CUFFT Code Examples

This section provides simple examples of 1D, 2D, and 3D complex and real data transforms that use the CUFFT to perform forward and inverse FFTs.

1D Complex-to-Complex Transforms

```
#define NX 256
#define BATCH 10

cufftHandle plan;
cufftComplex *data;
cudaMalloc((void**)&data, sizeof(cufftComplex)*NX*BATCH);

/* Create a 1D FFT plan. */
cufftPlan1d(&plan, NX, CUFFT_C2C, BATCH);

/* Use the CUFFT plan to transform the signal in place. */
cufftExecC2C(plan, data, data, CUFFT_FORWARD);

/* Inverse transform the signal in place. */
cufftExecC2C(plan, data, data, CUFFT_INVERSE);

/* Note:
(1) Divide by number of elements in data set to get back original data
(2) Identical pointers to input and output arrays implies in-place
    transformation
*/

/* Destroy the CUFFT plan. */
cufftDestroy(plan);
cudaFree(data);
```

1D Real-to-Complex Transforms

```
#define NX 256
#define BATCH 10

cufftHandle plan;
cufftComplex *data;
cudaMalloc((void**)&data, sizeof(cufftComplex)*(NX/2+1)*BATCH);

/* Create a 1D FFT plan. */
cufftPlan1d(&plan, NX, CUFFT_R2C, BATCH);

/* Use the CUFFT plan to transform the signal in place. */
cufftExecR2C(plan, (cufftReal*)data, data);

/* Destroy the CUFFT plan. */
cufftDestroy(plan);
cudaFree(data);
```

2D Complex-to-Complex Transforms

```
#define NX 256
#define NY 128

cufftHandle plan;
cufftComplex *idata, *odata;
cudaMalloc((void**)&idata, sizeof(cufftComplex)*NX*NY);
cudaMalloc((void**)&odata, sizeof(cufftComplex)*NX*NY);

/* Create a 2D FFT plan. */
cufftPlan2d(&plan, NX, NY, CUFFT_C2C);

/* Use the CUFFT plan to transform the signal out of place. */
cufftExecC2C(plan, idata, odata, CUFFT_FORWARD);

/* Note: idata != odata indicates an out-of-place transformation
to CUFFT at execution time. */
```

```
/* Inverse transform the signal in place */
cufftExecC2C(plan, odata, odata, CUFFT_INVERSE);

/* Destroy the CUFFT plan. */
cufftDestroy(plan);
cudaFree(idata); cudaFree(odata);
```

2D Complex-to-Real Transforms

```
#define NX 256
#define NY 128

cufftHandle plan;
cufftComplex *idata;
cufftReal *odata;
cudaMalloc((void**)&idata, sizeof(cufftComplex)*NX*NY);
cudaMalloc((void**)&odata, sizeof(cufftReal)*NX*NY);

/* Create a 2D FFT plan. */
cufftPlan2d(&plan, NX, NY, CUFFT_C2R);

/* Use the CUFFT plan to transform the signal out of place. */
cufftExecC2R(plan, idata, odata);

/* Destroy the CUFFT plan. */
cufftDestroy(plan);
cudaFree(idata); cudaFree(odata);
```

3D Complex-to-Complex Transforms

```
#define NX 64
#define NY 64
#define NZ 128

cufftHandle plan;
cufftComplex *data1, *data2;
cudaMalloc((void**)&data1, sizeof(cufftComplex)*NX*NY*NZ);
cudaMalloc((void**)&data2, sizeof(cufftComplex)*NX*NY*NZ);

/* Create a 3D FFT plan. */
cufftPlan3d(&plan, NX, NY, NZ, CUFFT_C2C);

/* Transform the first signal in place. */
cufftExecC2C(plan, data1, data1, CUFFT_FORWARD);

/* Transform the second signal using the same plan. */
cufftExecC2C(plan, data2, data2, CUFFT_FORWARD);

/* Destroy the CUFFT plan. */
cufftDestroy(plan);
cudaFree(data1); cudaFree(data2);
```
